


# Convolutional Neural Networks (CNN)

E. Milgo


# Introduction

- A Convolutional Neural Network (CNN) is a type of deep neural network primarily used for analyzing visual data.
- CNN can extract higher representations of the image.
- It is particularly effective for tasks like image classification, object detection, and segmentation.
- In classical image classification you define the image features. CNN takes the image's raw pixel data, trains the model and then extracts the features for better classification.

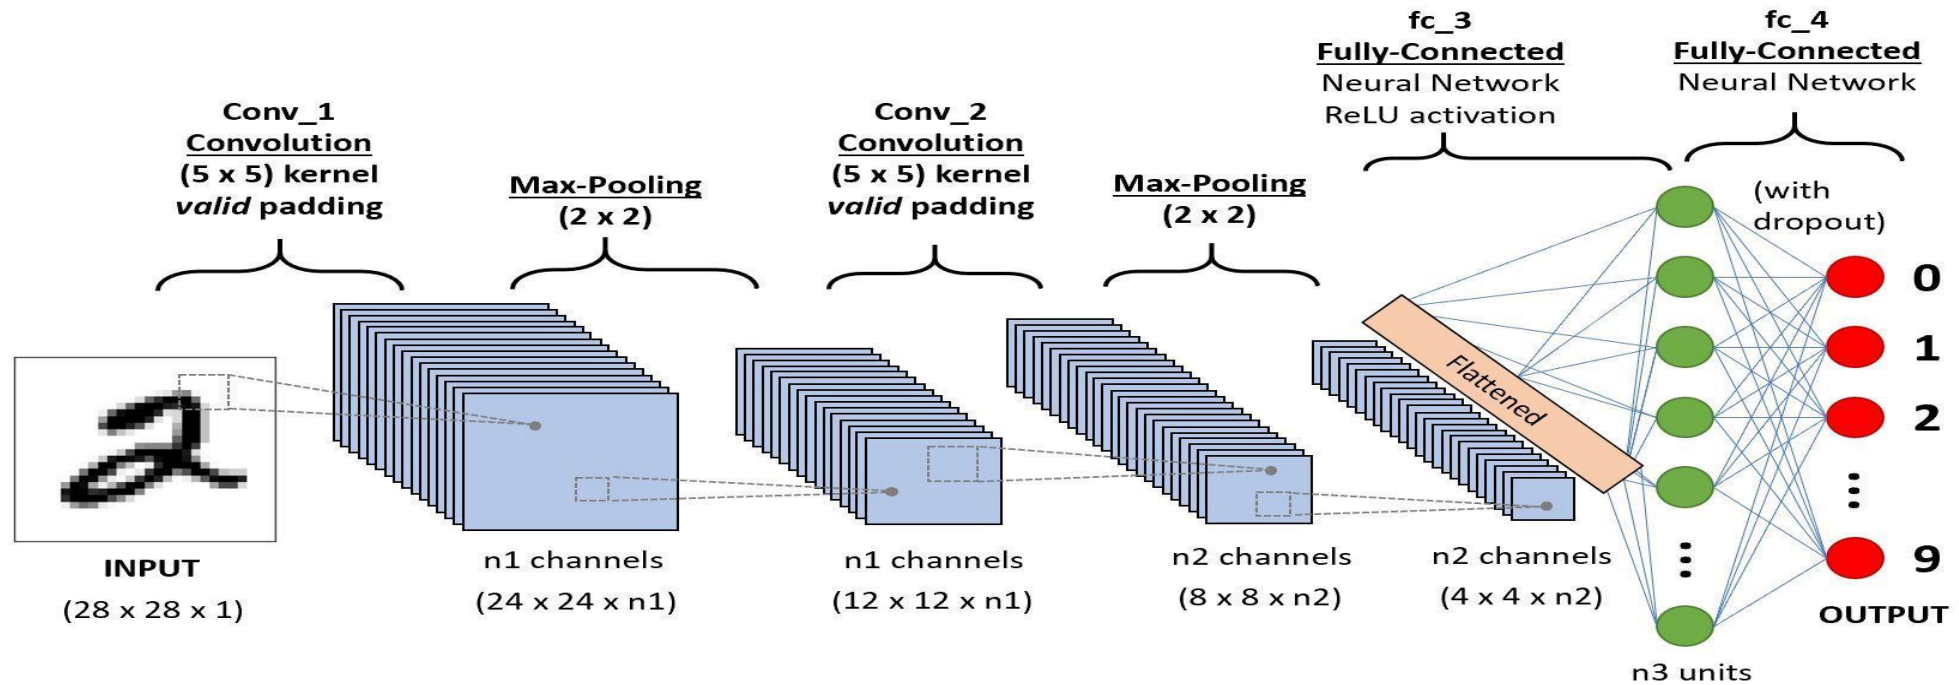
# Neural Network



# Convolutional Neural Networks (CNNs)



# How does it work





# Layers in CNN

- Input Layer: Image input (e.g., 28x28 pixels, grayscale or RGB).
- **Convolutional Layer** : Applies filters (kernels) to detect features like edges or textures.
- Activation Function (ReLU): Introduces non-linearity.
- **Pooling Layer (Max Pooling)**: Downsamples the image to reduce dimensions and computation.
- **Fully Connected Layers (Dense)** : Flattened features are passed through dense layers for classification.
- Output Layer: Final predictions

# Convolutional Layer

- A Convolutional Layer applies a set of filters (also called kernels) to the input (e.g., an image or feature map) to extract features like **edges**, **corners**, **textures**, or more complex patterns.

## How It Works

- A filter (usually a small matrix like 3x3 or 5x5) slides over the input data.
- At each position, it performs element-wise multiplication and sums the result → this is called the convolution operation.
- The output is a feature map (also called an activation map).

|   |   |   |   |   |   |
|---|---|---|---|---|---|
| 3 | 1 | 1 | 2 | 8 | 4 |
| 1 | 0 | 7 | 3 | 2 | 6 |
| 2 | 3 | 5 | 1 | 1 | 3 |
| 1 | 4 | 1 | 2 | 6 | 5 |
| 3 | 2 | 1 | 3 | 7 | 2 |
| 9 | 2 | 6 | 2 | 5 | 1 |

Original image 6x6

"Convolution"  
 $\times$

|   |   |    |
|---|---|----|
| 1 | 0 | -1 |
| 1 | 0 | -1 |
| 1 | 0 | -1 |

Filter 3x3

=

|     |     |  |  |
|-----|-----|--|--|
| -7  | ... |  |  |
| ... | ... |  |  |
|     |     |  |  |
|     |     |  |  |

Output 4x4

Result of the element-wise  
 product and sum of the  
 filter matrix and the original  
 image




# maxPooling Layer

- Max Pooling is a downsampling technique used in Convolutional Neural Networks (CNNs).
- It reduces the spatial dimensions (width and height) of the input feature map while retaining the most important information.

# How it works

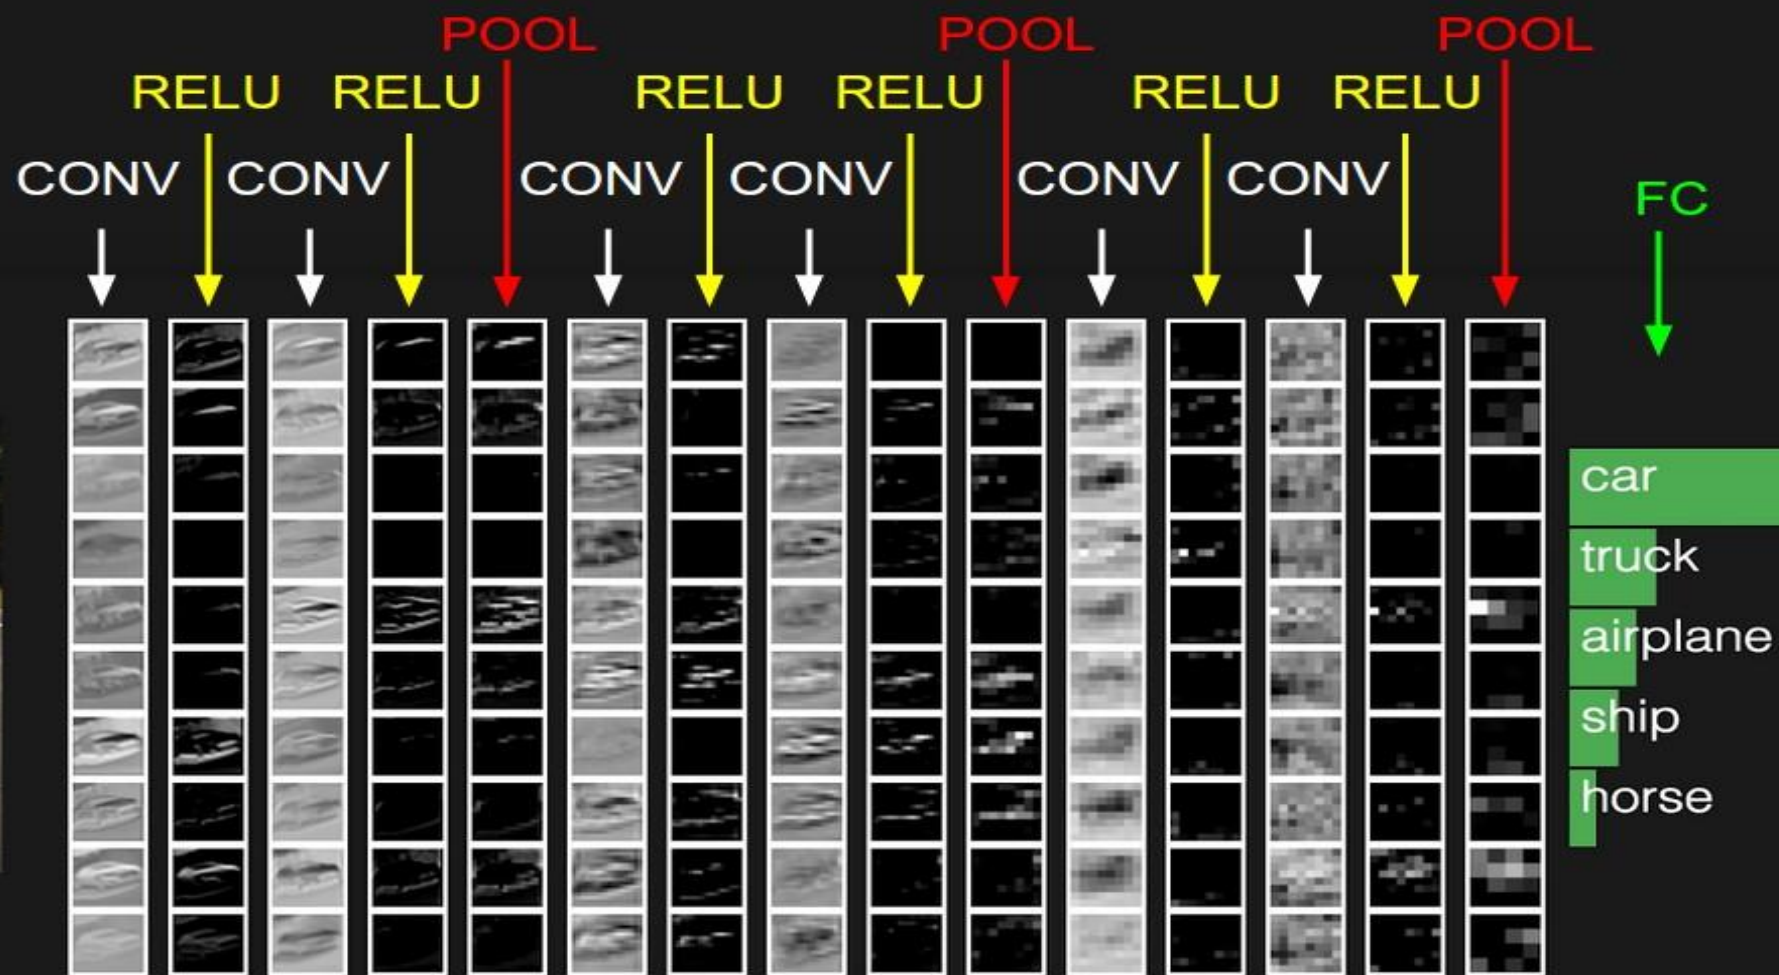
- A window (kernel) slides over the input feature map.
- For each window, it outputs the maximum value within that window.
- Common window size:  $2 \times 2$  with stride 2 (i.e., non-overlapping).

# MAX POOLING



# Fully Connected Layer (FC)

- A Fully Connected Layer (FC) is a standard neural network layer where every neuron is connected to every neuron in the previous layer.
- It's used at the end of the CNN, after the convolutional and pooling layers, to:
  - Flatten the spatial features
  - Interpret them
  - Produce the final output, like class scores





# Application of CNN

- Image classification (LeNet, InceptionNet, ResNet)
- Image segmentation (UNet, FCNN, RCNN)
- Image Generation(GAN)